

Saiakhil Chilaka

(984) 260-2526 | sc3484@cornell.edu | linkedin.com/in/akhilchilaka | github.com/akhilc08 | akhilchilaka.com

EDUCATION

Cornell University, College of Arts & Sciences

Ithaca, NY

B.A. in Computer Science

Expected May 2028

- **Relevant Coursework:** Distributed Systems for Machine Learning, Intro Machine Learning, Systems, Data Structures

EXPERIENCE

MotherDuck

May 2026 – August 2026

Software Engineering Intern

New York, NY

- Built the internal evaluation harness to benchmark AI agents across 40 MotherDuck tasks over the CLI & MCP
- Examined failed agent trajectories to write CLI guides for 3 features, taking task category accuracies from ~16% to 100%
- Reduced agent token cost on the MotherDuck CLI up to 57% through A/B testing 5 output design principles
- Designed & shipped CLI secret management commands, enabling AI agents deploy workflows without seeing credentials

Vanasse Hangen Brustlin, Inc.

January 2026 – May 2026

Software Engineering Intern

Remote/Raleigh, NC

- Designed an agentic PDF-to-form pipeline for car crash reports, cutting analyst time by 66% over 10k+ crashes annually
- Engineered an automated geocoder & RAG classification layer w/ project docs as context, achieving ~90% accuracy

Duke University

August 2024 – December 2024

Research Intern

Durham, NC

- Published first-author IEEE paper on CNN ensemble for Multiple Sclerosis diagnosis; 97.5% accuracy & 97.8% AUC-ROC
- Fine-tuned VGG19/ResNet50 layers via depth-wise separable convolutions, reducing model size 14x for edge deployment
- Injected Gaussian white noise alongside rotation, shear, and zoom augmentation to improve generalization by up to 7%

PROJECTS

Agentic Polymarket Predictor | *Python, Neo4j, Qdrant, Claude API, FastAPI, Modal*

- Built a 5-stage agentic predictor for Polymarket, achieving 30.5% out-of-sample ROI on 950+ markets and 140+ trades
- Architected a hybrid RAG pipeline over Neo4j & Qdrant via Reciprocal Rank Fusion with a BGE reranker
- Beat or matched the market on Brier scores on the 53% of markets the model chose, validated across multiple evaluations
- Caught 96% of hallucinated claims via a 4-dimension LLM judge pipeline to detect priming, recency, & confirmation bias

Pass@k is Not a Property of a Model | *Python, PyTorch, verl, AWS, SkyPilot*

- Published first-author preprint on evaluating RL-trained LLMs & methods for training/testing; submitted to conference
- Showed central results of *Limit of RLVR* (Yue et al.) fall 3x (36 to 12 pts) by altering the wording of the test questions
- Evaluated that an ensemble of 3 independently trained Qwen2.5 runs beats one model by 6% at matched sample budget
- Showed the main reasoning benchmark (MATH-500) is 98.8% solved by base models it tests, leaving 6/500 trials usable

First Person Boxing Game | *Python, PyTorch, ONNX, MediaPipe, Rust, TypeScript, Docker*

- Built a 2 player boxing game engine that tracks movement through your camera and simulates a boxing match, like a Wii
- Designed an ML model to classify punches into 5 types (ex. jab, hook) at 80% accuracy & compressed it to run in browser
- Architected the game engine using 33 tracked body points per player 60 times a second, works out from joint speed whether a hit landed and where, then runs damage with hit thresholds calibrated to each player's own reach and speed

TECHNICAL SKILLS

Languages: Python, TypeScript, R, SQL, Java

Frameworks & Libraries: PyTorch, TensorFlow, scikit-learn, LightGBM, verl, FastAPI, HuggingFace Transformers

Databases & Infrastructure: PostgreSQL, Redis, Supabase, Neo4j, Qdrant, Celery, Modal, AWS ECS, Docker, Git

ML/AI: MCP, RAG, Agent Evaluation, LoRA, Knowledge Distillation, Reinforcement Learning, Reranking, Vector Search

PUBLICATIONS

Chilaka et al. "Pass@k Is Not a Property of a Model," Preprints 2026, 2026081495, doi: 10.20944/preprints202608.1495.v1

Chilaka et al. "CNN Ensemble MS Lesions Diagnosis for Embedded Systems," *IEEE*, doi: 10.1109/MCSocC64144.2024.00079